

```
In [404... #There isnt enough data
import numpy as np
import pandas as pd

wine = pd.read_excel(r"C:\Users\djbro\OneDrive\Desktop\Multiple Linear Regression\Wine\W
wine1 = pd.read_excel(r"C:\Users\djbro\OneDrive\Desktop\Multiple Linear Regression\Wine\
```

```
In [405... #Concatenate wine and wine1
a=[wine,wine1]
wine = pd.concat(a)
wine
```

Out[405]:

	Year	Price	WinterRain	AGST	HarvestRain	Age	FrancePop
0	1979	6.9541	717	16.1667	122	4	54835.832
1	1980	6.4979	578	16.0000	74	3	55110.236
0	1952	7.4950	600	17.1167	160	31	43183.569
1	1953	8.0393	690	16.7333	80	30	43495.030
2	1955	7.6858	502	17.1500	130	28	44217.857
3	1957	6.9845	420	16.1333	110	26	45152.252
4	1958	6.7772	582	16.4167	187	25	45653.805
5	1959	8.0757	485	17.4833	187	24	46128.638
6	1960	6.5188	763	16.4167	290	23	46583.995
7	1961	8.4937	830	17.3333	38	22	47128.005
8	1962	7.3880	697	16.3000	52	21	48088.673
9	1963	6.7127	608	15.7167	155	20	48798.990
10	1964	7.3094	402	17.2667	96	19	49356.943
11	1965	6.2518	602	15.3667	267	18	49801.821
12	1966	7.7443	819	16.5333	86	17	50254.966
13	1967	6.8398	714	16.2333	118	16	50650.406
14	1968	6.2435	610	16.2000	292	15	51034.413
15	1969	6.3459	575	16.5500	244	14	51470.276
16	1970	7.5883	622	16.6667	89	13	51918.389
17	1971	7.1934	551	16.7667	112	12	52431.647
18	1972	6.2049	536	14.9833	158	11	52894.183
19	1973	6.6367	376	17.0667	123	10	53332.805
20	1974	6.2941	574	16.3000	184	9	53689.610
21	1975	7.2920	572	16.9500	171	8	53955.042
22	1976	7.1211	418	17.6500	247	7	54159.049
23	1977	6.2587	821	15.5833	87	6	54378.362
24	1978	7.1860	763	15.8167	51	5	54602.193

```
In [406... import seaborn as sns
```

```
#sns.pairplot(wine)
```

In [407...

```
# Checking for null values
print(wine.info())

# Checking for outliers
print(wine.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 27 entries, 0 to 24
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Year            27 non-null    int64
1   Price           27 non-null    float64
2   WinterRain      27 non-null    int64
3   AGST            27 non-null    float64
4   HarvestRain     27 non-null    int64
5   Age             27 non-null    int64
6   FrancePop       27 non-null    float64
dtypes: float64(3), int64(4)
memory usage: 1.7 KB
None
```

	Year	Price	WinterRain	AGST	HarvestRain	Age	\
count	27.000000	27.000000	27.000000	27.000000	27.000000	27.000000	
mean	1966.814815	7.041948	608.407407	16.477781	144.814815	16.185185	
std	8.246384	0.634590	129.034956	0.659189	73.065849	8.246384	
min	1952.000000	6.204900	376.000000	14.983300	38.000000	3.000000	
25%	1960.500000	6.508350	543.500000	16.150000	88.000000	9.500000	
50%	1967.000000	6.984500	600.000000	16.416700	123.000000	16.000000	
75%	1973.500000	7.441500	705.500000	17.008350	185.500000	22.500000	
max	1980.000000	8.493700	830.000000	17.650000	292.000000	31.000000	

	FrancePop
count	27.000000
mean	50085.443963
std	3792.998764
min	43183.569000
25%	46856.000000
50%	50650.406000
75%	53511.207500
max	55110.236000

In [408...

```
from sklearn.model_selection import train_test_split

# We specify random seed so that the train and test data set always have the same rows,
np.random.seed(0)
df_train, df_test = train_test_split(wine, train_size = 0.7, test_size = 0.3, random_sta
```

In [409...

```
#Re-scaling the Features
#We can see that all the columns have
#smaller integer values in the dataset
#except the area column. So it is important to
#re-scale the variables so that they all have a comparable scale.
#If we don't have relative scales, then some of the regression model
#coefficients will be of different units compared to the other coefficients.

#To do that, we use the MinMax scaling method.
```

In [410...

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()

# Applying scaler() to all the columns except the 'yes-no' and 'dummy' variables
num_vars = ['Year', 'Price', 'WinterRain', 'AGST', 'HarvestRain', 'Age', 'FrancePop']
df_train[num_vars] = scaler.fit_transform(df_train[num_vars])
```

df\_train

Out[410]:

	Year	Price	WinterRain	AGST	HarvestRain	Age	FrancePop
18	0.703704	0.000000	0.352423	0.00000	0.472441	0.296296	0.809211
14	0.555556	0.016865	0.515419	0.48668	1.000000	0.444444	0.649096
1	1.000000	0.128015	0.444934	0.40668	0.141732	0.000000	1.000000
8	0.333333	0.516908	0.707048	0.52668	0.055118	0.666667	0.395485
24	0.925926	0.428653	0.852423	0.33336	0.051181	0.074074	0.956261
23	0.888889	0.023506	0.980176	0.24000	0.192913	0.111111	0.936990
6	0.259259	0.137146	0.852423	0.57336	0.992126	0.740741	0.265941
4	0.185185	0.250044	0.453744	0.57336	0.586614	0.814815	0.185858
2	0.074074	0.647020	0.277533	0.86668	0.362205	0.925926	0.062231
16	0.629630	0.604422	0.541850	0.67336	0.200787	0.370370	0.725201
7	0.296296	1.000000	1.000000	0.94000	0.000000	0.703704	0.312777
5	0.222222	0.817372	0.240088	1.00000	0.586614	0.777778	0.226738
20	0.777778	0.038972	0.436123	0.52668	0.574803	0.222222	0.877693
1	0.000000	0.801468	0.691630	0.70000	0.165354	1.000000	0.000000
0	0.962963	0.327333	0.751101	0.47336	0.330709	0.037037	0.976375
19	0.740741	0.188658	0.000000	0.83336	0.334646	0.259259	0.846974
13	0.518519	0.277394	0.744493	0.50000	0.314961	0.481481	0.616035
10	0.407407	0.482567	0.057269	0.91336	0.228346	0.592593	0.504676

```
In [411... # Dividing the training data set into X and Y
y_train = df_train.pop('Price')
X_train = df_train
```

```
In [412... #Build a linear model

import statsmodels.api as sm
X_train_lm = sm.add_constant(X_train)

lr_1 = sm.OLS(y_train, X_train_lm).fit()

lr_1.summary()
```

```
C:\Users\djbro\anaconda3\lib\site-packages\scipy\stats\_stats_py.py:1772: UserWarning: kurtosistest only valid for n>=20 ... continuing anyway, n=18
warnings.warn("kurtosistest only valid for n>=20 ... continuing ")
```

Out[412]:

OLS Regression Results

<b>Dep. Variable:</b>	Price	<b>R-squared:</b>	0.855
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.795
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	14.17
<b>Date:</b>	Fri, 23 Dec 2022	<b>Prob (F-statistic):</b>	0.000111
<b>Time:</b>	18:09:13	<b>Log-Likelihood:</b>	13.551
<b>No. Observations:</b>	18	<b>AIC:</b>	-15.10

<b>Df Residuals:</b>	12	<b>BIC:</b>	-9.760			
<b>Df Model:</b>	5					
<b>Covariance Type:</b>	nonrobust					
	<b>coef</b>	<b>std err</b>	<b>t</b>	<b>P&gt; t </b>	<b>[0.025</b>	<b>0.975]</b>
<b>const</b>	-7.523e+12	1.07e+13	-0.706	0.494	-3.07e+13	1.57e+13
<b>Year</b>	7.523e+12	1.07e+13	0.706	0.494	-1.57e+13	3.07e+13
<b>WinterRain</b>	0.2413	0.157	1.537	0.150	-0.101	0.583
<b>AGST</b>	0.6482	0.186	3.486	0.004	0.243	1.053
<b>HarvestRain</b>	-0.5362	0.145	-3.709	0.003	-0.851	-0.221
<b>Age</b>	7.523e+12	1.07e+13	0.706	0.494	-1.57e+13	3.07e+13
<b>FrancePop</b>	-0.4490	0.902	-0.498	0.627	-2.413	1.515
<b>Omnibus:</b>	1.522	<b>Durbin-Watson:</b>	1.499			
<b>Prob(Omnibus):</b>	0.467	<b>Jarque-Bera (JB):</b>	0.899			
<b>Skew:</b>	0.091	<b>Prob(JB):</b>	0.638			
<b>Kurtosis:</b>	1.920	<b>Cond. No.</b>	9.09e+14			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 5.73e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
In [413... #Recursive Feature Elimination (RFE)
#RFE is an automatic process where we don't need to select
#variables manually. We follow the same steps we have done earlier
#until Re-scaling the features and dividing the data into X and Y.

#We will use the LinearRegression function from sklearn
#for RFE (which is a utility from sklearn)
```

```
In [414... # Importing RFE and LinearRegression
from sklearn.feature_selection import RFE
from sklearn.linear_model import LinearRegression
```

```
In [415... # Running RFE with the output number of the variable equal to 10
lm = LinearRegression()
lm.fit(X_train, y_train)

rfe = RFE(lm,n_features_to_select=10) # running RFE
rfe = rfe.fit(X_train, y_train)

list(zip(X_train.columns,rfe.support_,rfe.ranking_))
```

```
Out[415]: [('Year', True, 1),
('WinterRain', True, 1),
('AGST', True, 1),
('HarvestRain', True, 1),
('Age', True, 1),
('FrancePop', True, 1)]
```

In [416...

```

# Creating X_test dataframe with RFE selected variables
col = ['Year', 'WinterRain', 'AGST', 'HarvestRain', 'Age', 'FrancePop']
X_train_rfe = X_train[col]

# Adding a constant variable
import statsmodels.api as sm
X_train_rfe = sm.add_constant(X_train_rfe)

lm = sm.OLS(y_train, X_train_rfe).fit() # Running the linear model

print(lm.summary())

```

#### OLS Regression Results

```

=====
Dep. Variable:          Price      R-squared:            0.855
Model:                  OLS        Adj. R-squared:       0.795
Method:                 Least Squares    F-statistic:       14.17
Date:                  Fri, 23 Dec 2022    Prob (F-statistic): 0.000111
Time:                  18:09:13      Log-Likelihood:    13.551
No. Observations:      18          AIC:                -15.10
Df Residuals:          12          BIC:                -9.760
Df Model:               5
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-7.523e+12	1.07e+13	-0.706	0.494	-3.07e+13	1.57e+13
Year	7.523e+12	1.07e+13	0.706	0.494	-1.57e+13	3.07e+13
WinterRain	0.2413	0.157	1.537	0.150	-0.101	0.583
AGST	0.6482	0.186	3.486	0.004	0.243	1.053
HarvestRain	-0.5362	0.145	-3.709	0.003	-0.851	-0.221
Age	7.523e+12	1.07e+13	0.706	0.494	-1.57e+13	3.07e+13
FrancePop	-0.4490	0.902	-0.498	0.627	-2.413	1.515

```

=====
Omnibus:                1.522    Durbin-Watson:          1.499
Prob(Omnibus):           0.467    Jarque-Bera (JB):        0.899
Skew:                    0.091    Prob(JB):                0.638
Kurtosis:                1.920    Cond. No.                9.09e+14
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 5.73e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

C:\Users\djbro\anaconda3\lib\site-packages\scipy\stats\\_stats\_py.py:1772: UserWarning: kurtosistest only valid for n>=20 ... continuing anyway, n=18  
 warnings.warn("kurtosistest only valid for n>=20 ... continuing ")

In [417...

```

X_train_new = X_train_rfe.drop(["Year"], axis = 1)
#X_train_new = X_train_rfe
# Adding a constant variable
import statsmodels.api as sm
X_train_lm = sm.add_constant(X_train_new)

lm = sm.OLS(y_train, X_train_lm).fit() # Running the linear model
print(lm.summary())

```

#### OLS Regression Results

```

=====
Dep. Variable:          Price      R-squared:            0.848
Model:                  OLS        Adj. R-squared:       0.785
Method:                 Least Squares    F-statistic:       13.38
Date:                  Fri, 23 Dec 2022    Prob (F-statistic): 0.000147
Time:                  18:09:13      Log-Likelihood:    13.114
No. Observations:      18          AIC:                -14.23

```

```

Df Residuals:      12      BIC:      -8.886
Df Model:          5
Covariance Type:  nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
const          0.5737      1.005      0.571      0.579      -1.617      2.764
WinterRain     0.1900      0.143      1.332      0.208      -0.121      0.501
AGST           0.6233      0.187      3.332      0.006       0.216      1.031
HarvestRain   -0.4856      0.129     -3.774      0.003      -0.766     -0.205
Age           -0.2898      0.937     -0.309      0.762      -2.331      1.751
FrancePop     -0.6222      0.889     -0.700      0.497      -2.559      1.314
=====
Omnibus:                1.931   Durbin-Watson:                1.696
Prob(Omnibus):          0.381   Jarque-Bera (JB):        1.281
Skew:                   0.402   Prob(JB):                0.527
Kurtosis:              1.970   Cond. No.                74.0
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

C:\Users\djbro\anaconda3\lib\site-packages\scipy\stats\\_stats\_py.py:1772: UserWarning: kurtosistest only valid for n>=20 ... continuing anyway, n=18  
 warnings.warn("kurtosistest only valid for n>=20 ... continuing ")

```

In [418.. X_train_new = X_train_new.drop(["WinterRain"], axis = 1)

# Adding a constant variable
import statsmodels.api as sm
X_train_lm = sm.add_constant(X_train_new)

lm = sm.OLS(y_train,X_train_lm).fit()    # Running the linear model
print(lm.summary())

```

#### OLS Regression Results

```

=====
Dep. Variable:          Price      R-squared:                0.825
Model:                  OLS      Adj. R-squared:            0.772
Method:                 Least Squares      F-statistic:          15.37
Date:                  Fri, 23 Dec 2022      Prob (F-statistic):    7.52e-05
Time:                  18:09:13      Log-Likelihood:        11.873
No. Observations:      18      AIC:                  -13.75
Df Residuals:          13      BIC:                  -9.295
Df Model:               4
Covariance Type:       nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
const          1.2164      0.908      1.340      0.203      -0.745      3.178
AGST           0.5132      0.173      2.971      0.011       0.140      0.886
HarvestRain   -0.5278      0.128     -4.110      0.001      -0.805     -0.250
Age           -0.7155      0.906     -0.789      0.444      -2.673      1.243
FrancePop     -1.0687      0.847     -1.261      0.229      -2.899      0.762
=====
Omnibus:                4.628   Durbin-Watson:                1.770
Prob(Omnibus):          0.099   Jarque-Bera (JB):        1.452
Skew:                   0.088   Prob(JB):                0.484
Kurtosis:              1.620   Cond. No.                63.1
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

C:\Users\djbro\anaconda3\lib\site-packages\scipy\stats\\_stats\_py.py:1772: UserWarning: k

```
kurtosistest only valid for n>=20 ... continuing anyway, n=18
warnings.warn("kurtosistest only valid for n>=20 ... continuing ")
```

```
In [419... vif = pd.DataFrame()
X = X_train_new
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
print(vif)
```

	Features	VIF
0	const	684.789942
1	AGST	1.598020
2	HarvestRain	1.101181
3	Age	65.360862
4	FrancePop	64.139693

```
In [420... X_train_new = X_train_new.drop(["FrancePop"], axis = 1)

# Adding a constant variable
import statsmodels.api as sm
X_train_lm = sm.add_constant(X_train_new)

lm = sm.OLS(y_train,X_train_lm).fit() # Running the linear model
print(lm.summary())
```

```
C:\Users\djbro\anaconda3\lib\site-packages\scipy\stats\_stats_py.py:1772: UserWarning: k
urtosistest only valid for n>=20 ... continuing anyway, n=18
warnings.warn("kurtosistest only valid for n>=20 ... continuing ")
```

#### OLS Regression Results

```
=====
Dep. Variable:          Price      R-squared:                0.804
Model:                  OLS       Adj. R-squared:            0.762
Method:                 Least Squares   F-statistic:          19.16
Date:                   Fri, 23 Dec 2022   Prob (F-statistic):    3.17e-05
Time:                   18:09:13    Log-Likelihood:        10.834
No. Observations:       18          AIC:                   -13.67
Df Residuals:           14          BIC:                   -10.11
Df Model:                3
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	0.0782	0.103	0.762	0.459	-0.142	0.298
AGST	0.5103	0.176	2.894	0.012	0.132	0.888
HarvestRain	-0.5511	0.130	-4.249	0.001	-0.829	-0.273
Age	0.4132	0.147	2.813	0.014	0.098	0.728

```
=====
Omnibus:                 1.570    Durbin-Watson:              1.872
Prob(Omnibus):            0.456    Jarque-Bera (JB):        0.917
Skew:                    0.107    Prob(JB):                0.632
Kurtosis:                 1.915    Cond. No.                8.22
=====
```

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [421... vif = pd.DataFrame()
X = X_train_new
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
print(vif)
```

	Features	VIF
0	const	8.392431
1	AGST	1.597728

```
2 HarvestRain 1.078253
3 Age 1.646593
```

In [422...

```
X_train_new = X_train_new.drop(["const"], axis = 1)

# Adding a constant variable
import statsmodels.api as sm
X_train_lm = sm.add_constant(X_train_new)

lm = sm.OLS(y_train,X_train_lm).fit() # Running the linear model
print(lm.summary())

vif = pd.DataFrame()
X = X_train_new
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
print(vif)
```

#### OLS Regression Results

```
=====
Dep. Variable: Price R-squared: 0.804
Model: OLS Adj. R-squared: 0.762
Method: Least Squares F-statistic: 19.16
Date: Fri, 23 Dec 2022 Prob (F-statistic): 3.17e-05
Time: 18:09:13 Log-Likelihood: 10.834
No. Observations: 18 AIC: -13.67
Df Residuals: 14 BIC: -10.11
Df Model: 3
Covariance Type: nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	0.0782	0.103	0.762	0.459	-0.142	0.298
AGST	0.5103	0.176	2.894	0.012	0.132	0.888
HarvestRain	-0.5511	0.130	-4.249	0.001	-0.829	-0.273
Age	0.4132	0.147	2.813	0.014	0.098	0.728

```
=====
Omnibus: 1.570 Durbin-Watson: 1.872
Prob(Omnibus): 0.456 Jarque-Bera (JB): 0.917
Skew: 0.107 Prob(JB): 0.632
Kurtosis: 1.915 Cond. No. 8.22
=====
```

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

	Features	VIF
0	AGST	5.317162
1	HarvestRain	2.188327
2	Age	5.476022

```
C:\Users\djbro\anaconda3\lib\site-packages\scipy\stats\_stats_py.py:1772: UserWarning: kurtosistest only valid for n>=20 ... continuing anyway, n=18
warnings.warn("kurtosistest only valid for n>=20 ... continuing ")
```

In [423...

```
#Since the p-values and VIF are in the desired range, we'll move forward with the analysis
```

In [424...

```
#The next step is the residual analysis of error terms.
```

```
#Residual Analysis
```

```
#So, let's check if the error terms are also normally distributed using a histogram.
```

In [425...


```
y_train_price = lm.predict(X_train_lm)
# Importing the required libraries for plots.
import matplotlib.pyplot as plt
import seaborn as sns
```



```
%matplotlib inline

# Plot the histogram of the error terms
fig = plt.figure()
sns.histplot((y_train - y_train_price), bins = 5)
```

Out[425]: <AxesSubplot:ylabel='Count'>



```
In [426... num_vars = ['Year', 'Price', 'WinterRain', 'AGST', 'HarvestRain', 'Age', 'FrancePop']
df_test[num_vars] = scaler.transform(df_test[num_vars])

y_test = df_test.pop('Price')
X_test = df_test

# Now let's use our model to make predictions.

# Creating X_test_new dataframe by dropping variables from X_test
X_test_new = X_test[X_train_new.columns]
# Adding a constant variable
X_test_new = sm.add_constant(X_test_new)

# Making predictions
y_pred = lm.predict(X_test_new)
```

```
In [427... from sklearn.metrics import r2_score
r2_score(y_true = y_test, y_pred = y_pred)
```

Out[427]: 0.5760134554123846

```
In [428... from sklearn import metrics
print(metrics.mean_absolute_error(y_test, y_pred))
print(metrics.mean_squared_error(y_test, y_pred))
print(np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
```

```
0.11568220627300144
0.01802667325935115
0.13426344721982655
```

In [429...

```
#The  $R^2$  value for the test data = 0.6481740917926483,  
#which is pretty similar to the train data.
```

```
#Since the  $R^2$  values for both the train and  
#test data are almost equal, the model we built is the best-fitted model.
```